ANTS: Adversarial Agentic Task Synthesis

Agentic Physical AI Framework for Industrial Automation: End-to-End Robot Policy Development, Hardening and Deployment

Team project with Daniel Totolici · Submitted to the Meta Global AI Hackathon

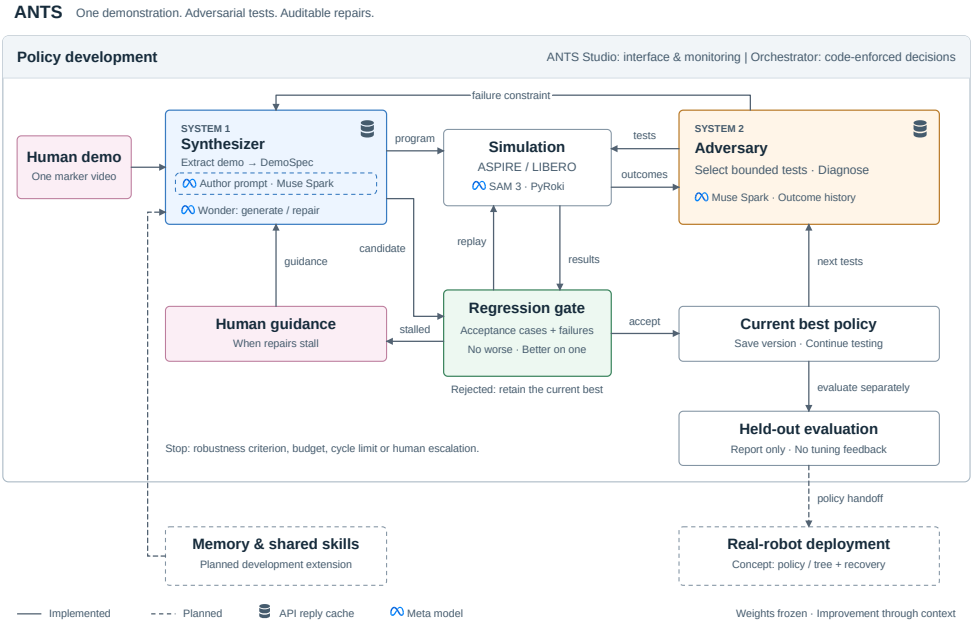
Code as policy | Muse Wonder | Muse Spark | SAM 3/3.1 | LIBERO-Pro

One Demo In, a Hardened Robot Program Out

SUBMITTED

Show a robot a task **once on video**. A frozen model writes the robot program. A second system attacks it until the edge cases are covered. **Nothing is trained.** The skill improves through context only.

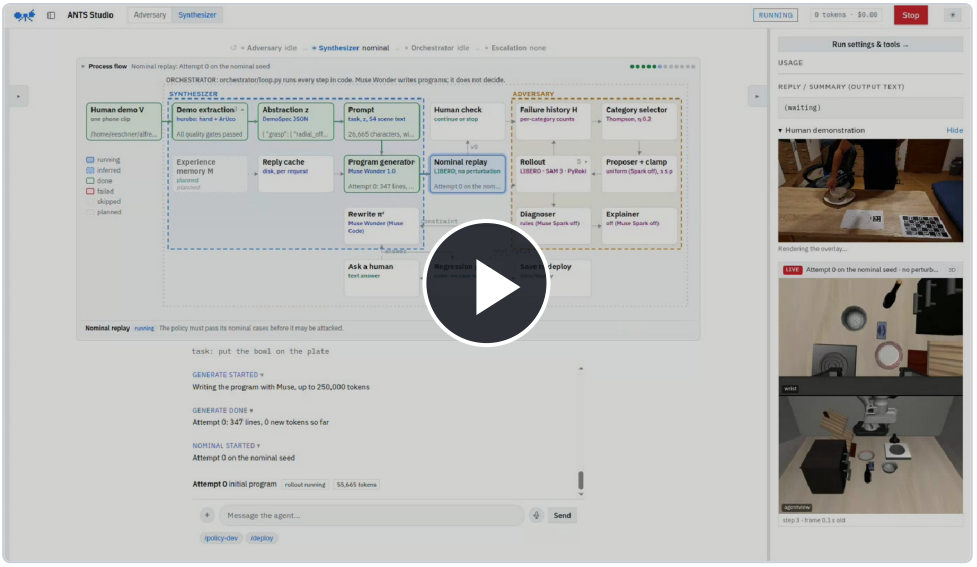
Submitted to the Meta Global AI Hackathon, theme *Push the Frontier*: use Meta's models to tackle challenges in robotics and beyond. ANTS pushes it with Meta's own models in both systems: Muse Wonder writes the robot program, Muse Spark plans and diagnoses the attacks, and SAM 3/3.1 gives the robot its perception.



Two systems, one goal: a robust policy. Solid boxes are built, dashed boxes are planned.

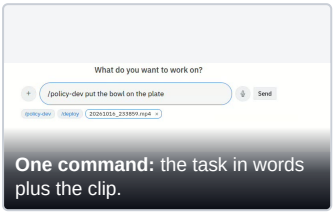
ANTS Studio: One Run, End to End

ANTS Studio is the console for the whole loop. One chat command starts a run from a phone clip; every stage shows its status live, the generated program runs in simulation next to the human demo, and the run waits for a person before the adversary attacks it.

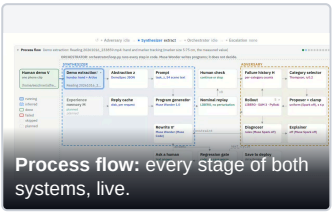


▶ Watch the 3-minute ANTS Studio walkthrough video: alfredowh.github.io/projects/ants.html#studio

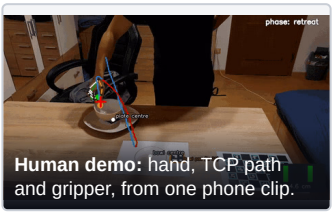
SHOW ONCE: SYSTEM 1



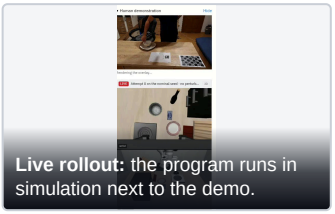
One command: the task in words plus the clip.



Process flow: every stage of both systems, live.

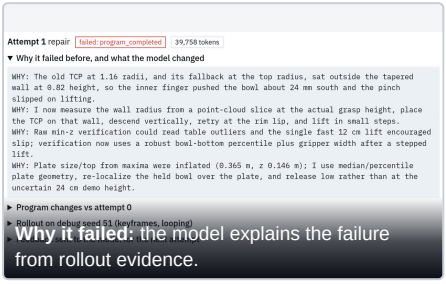


Human demo: hand, TCP path and gripper, from one phone clip.

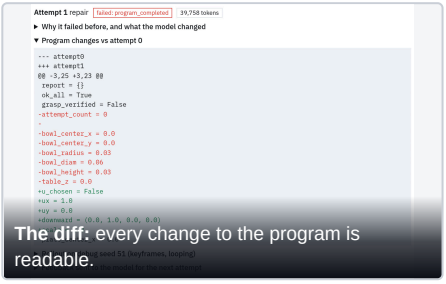


Live rollout: the program runs in simulation next to the demo.

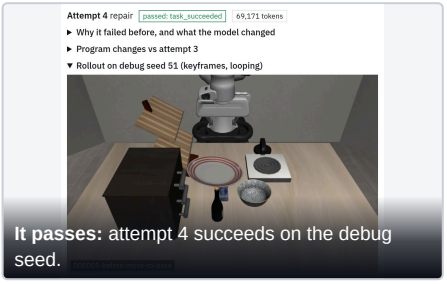
REPAIR LOOP: ONE LOGGED RUN, FOUR REPAIRS



Why it failed: the model explains the failure from rollout evidence.

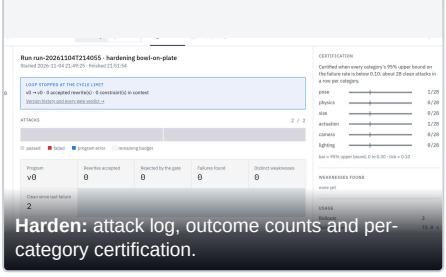


The diff: every change to the program is readable.

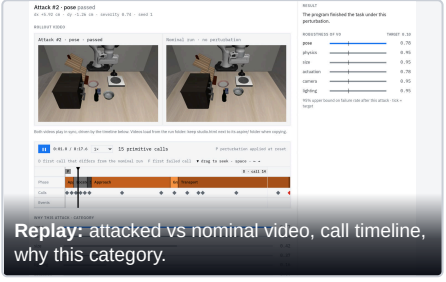


It passes: attempt 4 succeeds on the debug seed.

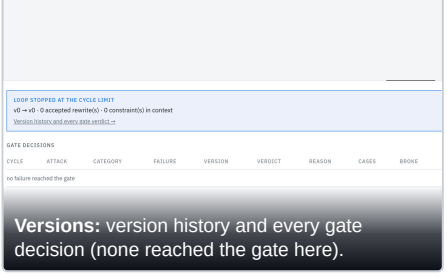
ADVERSARY TAB: SYSTEM 2



Harden: attack log, outcome counts and per-category certification.



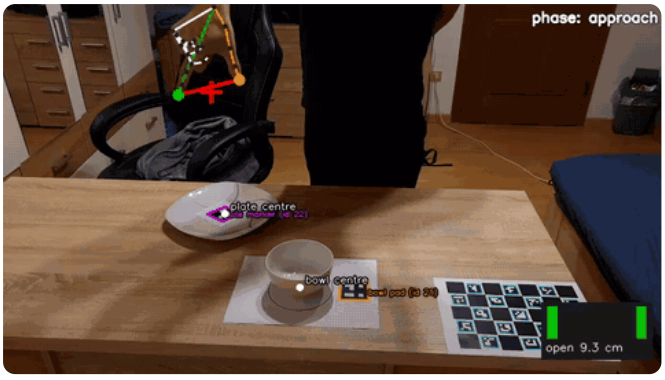
Replay: attacked vs nominal video, call timeline, why this category.



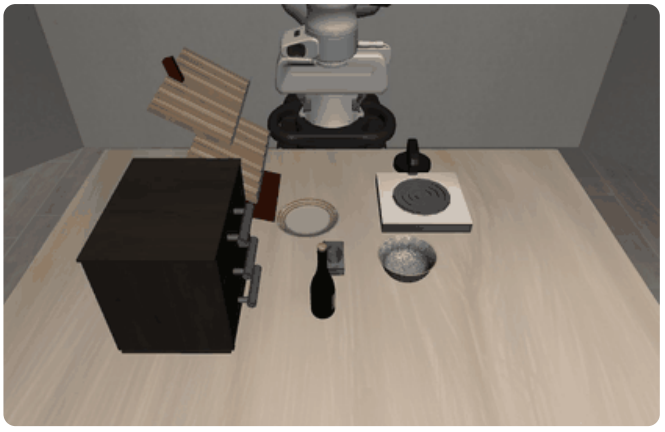
Versions: version history and every gate decision (none reached the gate here).

A Human Video as the Prompt

A text prompt says *"put the bowl on the plate"*. It cannot say *where* to grasp, *how high* to carry or *where* to release. A human video can. We track the hand and objects in one demo and turn them into numbers relative to object size: grasp the rim, approach from above, carry height, release over the plate. The VLM writes the robot program from these numbers instead of guessing them.



One human demo, with what the tracker reads.

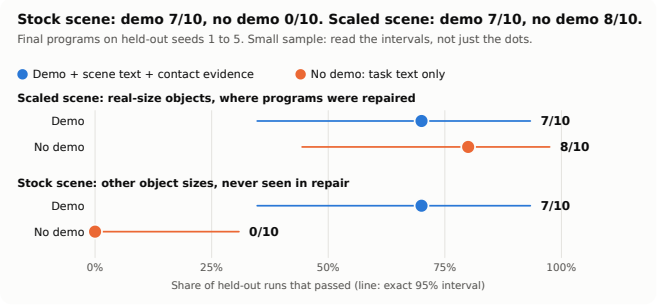


v0, the first candidate policy generated and tested by System 1, run in LIBERO-Pro.

Human Video Prompt vs. Text-Only Prompt

Human video prompt

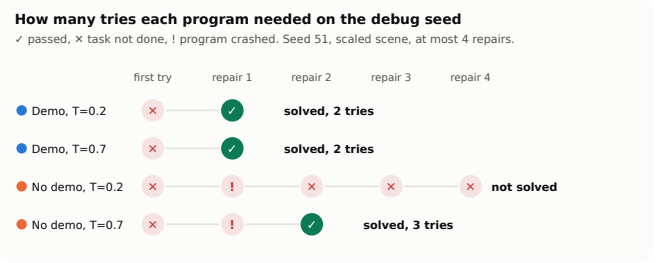
"Put the bowl on the plate" plus one tracked human demo. The numbers (grasp point, approach direction, carry height, release spot) come from the video, relative to object size.



Held-out pass rate with 95% intervals. Programs were repaired on the scaled objects only.

Text-only prompt (no demo)

"Put the bowl on the plate" and nothing else. The VLM has to guess where to grasp, how high to carry and where to release.



Tries needed on the debug seed (scaled objects).

Both arms were repaired on the scaled objects. There, text only catches up (8/10 vs 7/10), but on the debug seed the demo programs solved it in 2 tries while one text-only program never did. On other object sizes the text-only program fails (0/10), while the demo program holds (7/10): the video gives relative numbers that transfer across sizes. This is still an early hint, with one task, one clip and five seeds. The demo can also come from teleoperation, not only from video.

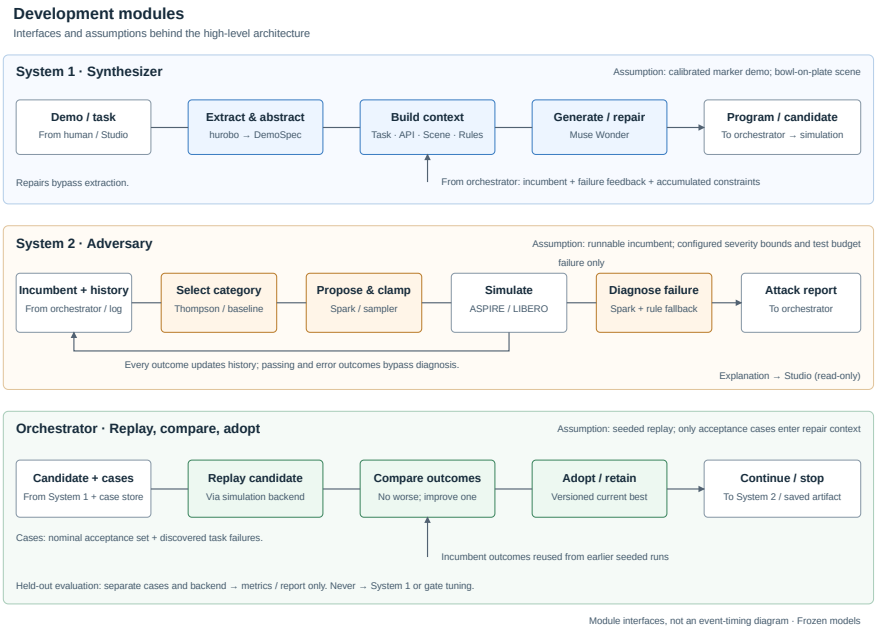
System 1 and System 2, Working on One Objective

Two systems run independently, talk through a small contract and share one goal: a policy that does not break.

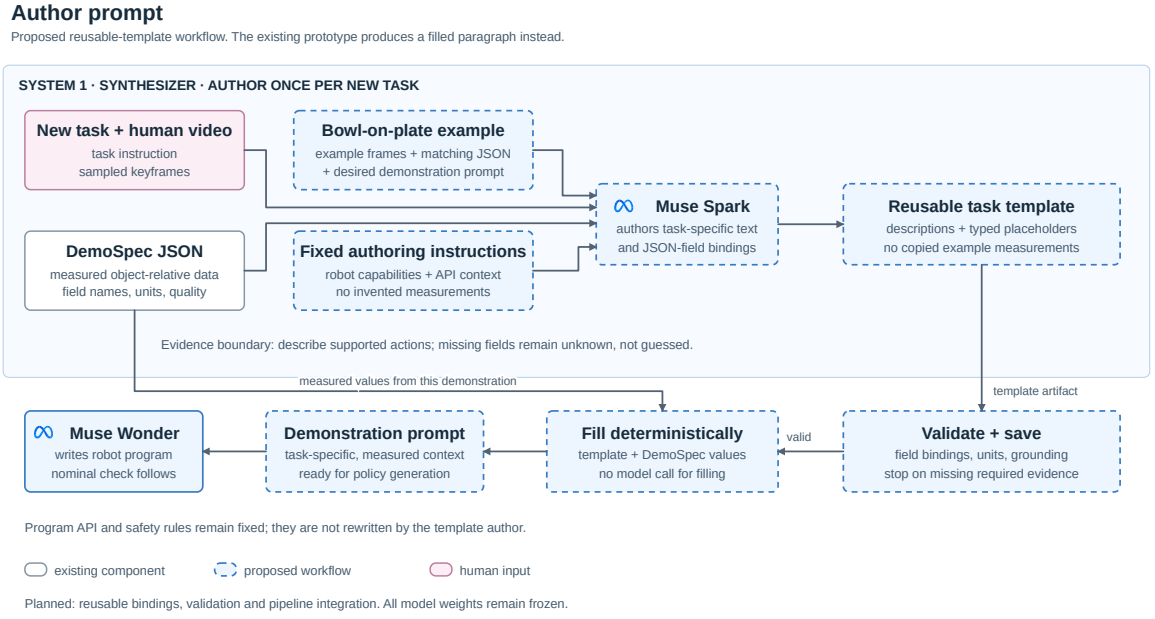
- **System 1, policy generator (synthesizer):** turns the prompt (human video, teleop or text) into a **program as policy**, checks it in simulation and repairs it.
- **System 2, failure case generator (adversary):** looks for the cases most likely to break the policy. It picks the weakest category, proposes perturbations within a bounded severity (pose, size, physics, actuation, camera, lighting), runs them and diagnoses each failure into a rule.
- **Gate:** every rule goes back into System 1's context. A rewrite is accepted only if it is no worse on all cases seen so far and better on one.

Each system makes the other better: the generator hardens the program, the adversary keeps finding what the hardening missed. When the loop stalls it asks a human, and the answer becomes a rule.

Note: I worked on the synthesizer and the demo-to-program idea. I did not look into the technical details of the adversary, which is Daniel's part.

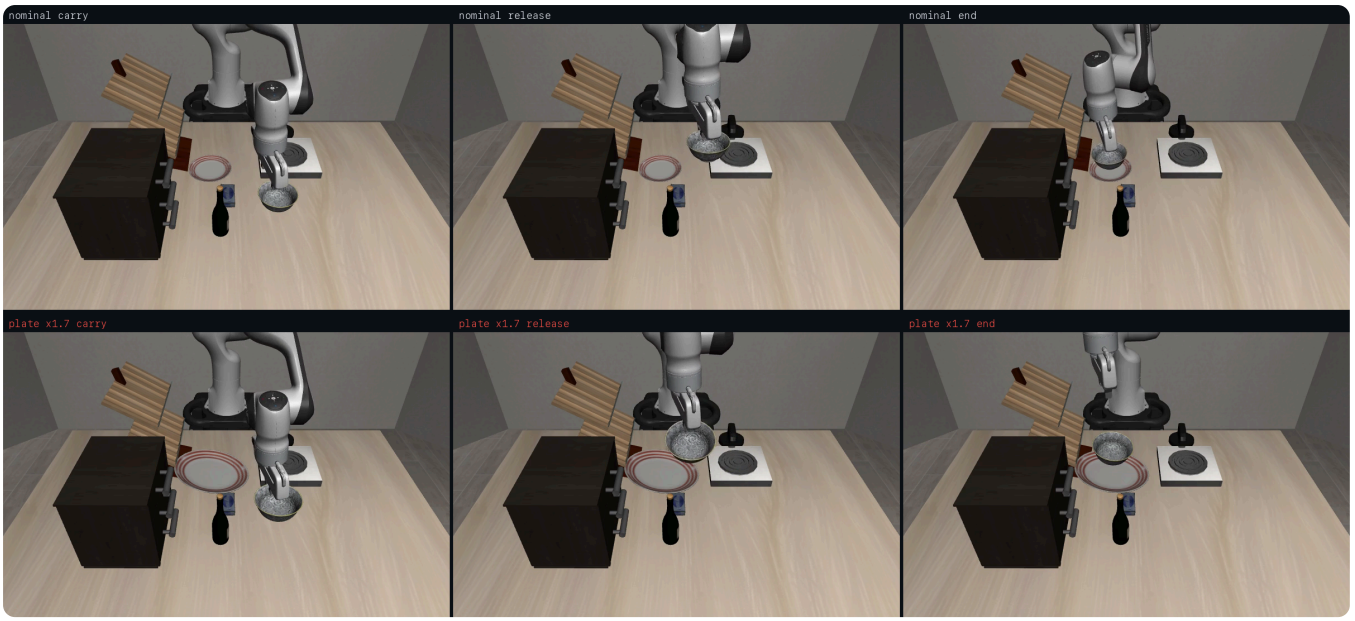


Modules and interfaces of the development loop.



Planned: an author step turns the demo evidence into a reusable task prompt before the program is written.

What the Adversary Found



Nominal (top) versus plate ×1.7 (bottom), same seed.

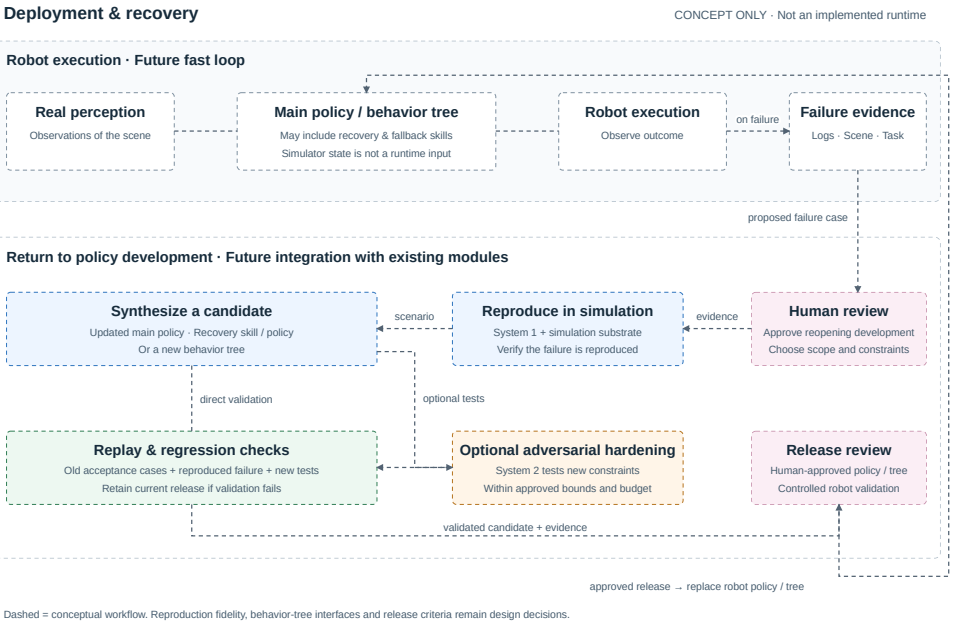
v0 is the first candidate policy generated and tested by System 1. Each accepted rewrite becomes the next version (v1, v2, ...). With a larger plate, SAM 3 measured the plate 11 cm tall instead of 2 cm. The program released the bowl about 9 cm too high. Muse Spark traced it to one line, `plate_top_z = max(...)`, Muse Wonder rewrote it and the gate accepted the fix.

Question	Result	From
One demo gives a working program?	v0 passes 47/50 held-out perturbed cases	759 LIBERO rollouts
Adversary finds a real weak spot?	0 failures in 599 rollouts with correct perception, 11 in 26 with a mis-read plate	one blind spot
Gate blocks bad fixes?	3 of 3 bad rewrites rejected	logged runs
Fix generalises?	**Not yet.** Held-out size cases: v0 49/50, v1 47/50	gate saw only 4 cases
Adaptive beats uniform testing?	**Not answered.** 2/200 vs 1/200 failures	v0 fails too rarely

Held-out cases use scene seeds the system never tuned on. Cost: about $12 of model tokens for the whole study. Limits are stated, not hidden.

Add-on: Recovery Policy

A single policy that handles every edge case becomes long, hard to read and hard to reuse. The architecture can instead let System 1 write a **main policy plus small recovery policies**. This works in the development phase and later in deployment: a production failure is reproduced in simulation, the adversary hardens the fix, and the new recovery skill is released after validation. A final step could compose main and recovery skills into a **behavior tree**, which keeps each part modular and reusable. Not built yet.



Deployment and recovery concept (not implemented).

Outlook

- **Comprehensive evaluation** of demo versus no demo, and a broader adversary, across more tasks and acceptance sets. Current tests are a start, not comprehensive.
- **Better planning:** cuRobo trajectory optimisation against a static 3D map of the scene. A scene reconstruction probe (TSDF to ESDF) already works.
- **Sim-to-real gap:** today the gap is assumed to be zero, so robust in simulation means robust in reality. Next: gap-aware test cases, or a real2sim digital twin that replicates the real scene 1:1.
- **Deployment phase:** today ANTS covers policy development only.
- **ANTS Studio demo: coming soon!** One window to watch the whole loop: the synthesizer chat with live cameras, and the adversary's attacks, replays and versions.
- **Memory:** store hardened policies and share skills across tasks.

Team

Alfredo: synthesizer and demo prompt. Daniel: adversary. Alfredo Wijaya Hardjoprawiro and Daniel Totolici. Demonstrations come from hurobo, which turns phone videos of a human hand into robot demos.

Alfredo Wijaya Hardjoprawiro — M.Sc. Artificial Intelligence, TH Ingolstadt

alh7017@thi.de alfredowijaya@gmail.com GitHub LinkedIn alfredowh.github.io